

King Fahd University of Petroleum & Minerals
College of Computing and Mathematics
Information and Computer Science Department
ICS 381: Principles of AI – First Semester 2022-2023 (221)
PA 2 – Programming Instructions

General Helpful Tips:

- You can submit your code on Gradescope any time before the submission deadline. The autograder will test your code and give you scores feedback on various test cases. These tests are to ease your mind that your implementation looks correct. However, note that more tests will be run on your implementation **after** the submission deadline. This is done to ensure that students do **not** base their implementations on passing the initial tests; i.e. the implementation should be based on the specifications in this document.
- **It is important that you adhere to the specifications and naming given in this document. Otherwise, the autograder will fail to run your implementation.**
- **Do not copy others' work. You can discuss general approaches with students, but do not direct or share coding solutions.**
- Submit the required files only: [[problem.py](#), [node.py](#), [search_algorithms.py](#)]
- Only import as minimally as possible. Do not use packages like numpy, pandas, etc unless otherwise stated. The autograder will check this and deduct points if extra packages are imported.
- **NOTE:** It is highly recommended that you test and submit your code after you implement each task chunk. It will be easier to debug your code. Each task typically builds on the previous task. Do not wait until you implement everything to test your code.

problem.py: Implement Problem Class

Your task is to implement a base class for search problems that we will use later to define other problems via inheritance. Specifically, implement a class with the following properties:

Class name	Inherits from
Problem	object

Problem Constructor	
Constructor arguments	Constructor body
initial_state, goal_state=None	Add arguments to self.

Problem Functions			
Name	Arguments	Returns	implementation
actions	state	Returns actions available in state	Since Problem is a base class, this function should just raise NotImplementedError.
result	state, action	Returns the transition-state that results from executing action in state.	Since Problem is a base class, this function should just raise NotImplementedError.
is_goal	state	Returns true if the given state is equal to the goal, otherwise false. Assume == is enough to check for equality.	Returns true if the given state is equal to the goal, otherwise false. Use == for equality.
action_cost	state1, action, state2	Returns the cost of transition from state1 to state2 using action.	For this base class, should return 1.
h	node	Returns the heuristic value of the given node.	For this base class, should return 0.

[node.py](#): Implement Node Class

Your task is to implement a Node class that we will use as the building blocks for the search tree. Implement as follows:

Class name	Inherits from
Node	object

Node Constructor	
Constructor arguments	Constructor body
state, parent_node=None, action_from_parent=None, path_cost=0	Add arguments to self. In addition, add self.depth to be 0 if no parent-node or depth of the parent-node + 1.

Node class has only a single function:

```
def __lt__(self, other):  
    return self.state < other.state
```

This is just a safe-guard to break ties in the priority-queue using the state. See next page for priority queue class.

[search_algorithms.py](#): Implement search helper functions

Firstly, add the following implementation of PriorityQueue that we will use later for search algorithms:

```
import heapq

class PriorityQueue:
    def __init__(self, items=(), priority_function=(lambda x: x)):
        self.priority_function = priority_function
        self.pqueue = []
        # add the items to the PQ
        for item in items:
            self.add(item)

    """
    Add item to PQ with priority-value given by call to priority_function
    """
    def add(self, item):
        pair = (self.priority_function(item), item)
        heapq.heappush(self.pqueue, pair)

    """
    pop and return item from PQ with min priority-value
    """
    def pop(self):
        return heapq.heappop(self.pqueue)[1]

    """
    gets number of items in PQ
    """
    def __len__(self):
        return len(self.pqueue)
```

Now implement the following search helper functions in `search_algorithms.py`. Be sure to add imports for `problem.py` and `node.py`. Implement the following functions (note these functions are not part of a class, just regular python functions):

Name	Arguments	Returns	implementation
<code>expand</code>	<code>problem,</code> <code>node</code>	Returns the children nodes of the given node.	See figure 3.7 pg 91 in AIMI 4/E.
<code>get_path_actions</code>	<code>node</code>	Returns a list of actions to get to the given node. If node is None return empty list []. If parent of node is None return an empty list [].	Hint: follow the parent pointers to the root. Be sure to return in the correct order (not reverse order).
<code>get_path_states</code>	<code>node</code>	Returns a list of states to get to the given node. If node is None return empty list []	Hint: follow the parent pointers to the root. Be sure to return in the correct order (not reverse order).

[search_algorithms.py](#): Implement search algorithms

Your task is to implement search algorithms in `search_algorithms.py`. To make your job easier, first implement the best-first-search template. Then implement the main search algorithms as functions that call the best-first-search template with appropriate evaluation functions. Be sure to add imports for `problem.py` and `node.py`.

The following are the functions you need to implement:

Name	Arguments	Returns	implementation
<code>best_first_search</code>	<code>problem, f</code>	Returns the goal-node if found, otherwise returns <code>None</code> in case of failure .	See figure 3.7 pg 91 in AIMI 4/E. You can implement <i>reached</i> as a python dictionary with states as keys and nodes as values. Note that the argument <code>f</code> is a function (in python you can pass functions as arguments).
<code>best_first_search_treelike</code>	<code>problem, f</code>	Same as above.	Implement tree-like version of <code>best_first_search</code> . See slides
<code>breadth_first_search</code>	<code>problem, treelike=false</code>	Same as above.	Hint: implement as a call to <code>best_first_search</code> (or <code>treelike</code> version) with appropriate <code>f</code> function. Use the <code>treelike</code> flag to decide which version to call.
<code>depth_first_search</code>	<code>problem, treelike=false</code>	Same as above.	Same as above.
<code>uniform_cost_search</code>	<code>problem, treelike=false</code>	Same as above.	Same as above.
<code>greedy_search</code>	<code>problem, h, treelike=false</code>	Same as above.	Same as above. Note that the argument <code>h</code> is a heuristic function.
<code>astar_search</code>	<code>problem, h, treelike=false</code>	Same as above.	Same as above.


Feel free to add other functions to help you implement these algorithms. Hint: you can implement various functions to use as the `f` argument for certain search algorithms.

problem.py: Implement RouteProblem Class

Your task is to implement a class for route problems in problem.py. To ease your implementation, be sure to inherit the Problem class and override its functions.

Class name		Inherits from	
RouteProblem		Problem	
RouteProblem Constructor			
Constructor arguments		Constructor body	
initial_state, goal_state=None, map_graph=None, map_coords=None		Add arguments to self. Assume map_graph is a dictionary where keys are (v1, v2) tuple indicating a directed edge between v1 and v2 states, values give the cost between v1 and v2. Assume map_coords is a dictionary where keys are states, and values are (x, y) tuple indicating the xy-coordinates in 2D.	
RouteProblem Functions to Override			
Name	Arguments	Returns	implementation
actions	state	Returns the states that can be reached by state If no state is reachable, then return an empty list []	Hint: although not necessary, consider making a self.neighbors dictionary in the constructor to help you.
result	state, action	Returns the transition-state that results from executing action in state. If no transition is possible, then return state (i.e. do not move)	Assume action will be a state. For example, result('A', 'B') is the action of going from A to B.
action_cost	state1, action, state2	Returns the cost of transition from state1 to state2 using action.	Hint: use map_graph Assume action == state2.
h	node	Returns the Euclidean distance to the goal.	Hint: use map_coords Note: make sure to return 0 if node is a goal-node.

To clarify the usage of RouteProblem, consider the following map of cities.



Assume that start state is A, and the goal is E, assume all costs are 1. Then you can construct this graph as a dictionary:

```
example_map_graph = {  
    ('A', 'B'): 1,  
    ('A', 'C'): 1,  
    ('A', 'D'): 1,  
    ('B', 'A'): 1,  
    ('B', 'C'): 1,  
    ('B', 'E'): 1,  
    ('C', 'B'): 1  
}  
  
example_coords = {  
    'A': (1,2),  
    'B': (0,1),  
    'C': (1,1),  
    'D': (2,1),  
    'E': (0,0),  
}
```


Then create a RouteProblem object as:

```
example_route_problem = RouteProblem(initial_state='A', goal_state='E',  
                                     map_graph=example_map_graph,  
                                     map_coords=example_coords)  
  
goal_node = breadth_first_search(example_route_problem, tree_like=False)
```

You can try testing your search algorithms by passing a route-problem. Try testing the practice problem from the slides.

problem.py: Implement GridProblem Class

Your task is to implement a class for grid problems in `problem.py`. To ease your implementation, be sure to inherit the `Problem` class and override its functions. Unlike route-problems where we pass the entire graph in the constructor, in a grid problem, we just need certain static information about the environment.

A grid problem instance consists of an N-by-M grid (N rows and M columns). Some grid locations contain walls that the agent cannot move into; we will use a list of locations `wall_coords` to store these wall locations. Some grid locations contain food that the agent can consume `food_coords`. The agent starts at some grid location (xA, yA). The agent can move left, right, up, and down one grid. The agent cannot move outside the bounds of the N-by-M grid. The goal is to consume all food on the grid.


With this formulation, we need to keep track of which food was eaten. So, a `state` is a tuple of two elements: the current location (xA, yA) of the agent, and a Boolean **tuple** `food_eaten` of the same size as `food_coords` where `food_eaten[i]` indicates if `food_coords[i]` was eaten (true) or not (false). Initially, this tuple should be set to false (all food not eaten at start).

Note: the reason we represent `food_eaten` as a tuple instead of a list is because python dictionaries do not allow keys to be lists. Recall that we want to use a reached dictionary for our graph-like search algorithms. So, our states cannot have list objects. Moreover, you cannot modify or change tuples, they are immutable. However, you can convert a tuple X to list as `list(X)`, perform modifications, then convert back to tuple as `tuple(mylist)`. This may be convenient when you create the transition-states in the `result` function.

Class name		Inherits from	
GridProblem		Problem	
GridProblem Constructor			
Constructor arguments		Constructor body	
initial_state, N, M, wall_coords, food_coords		<p>Add arguments to self. You can assume the following: initial_state is a pair (xA, yA) indicating the agent's starting grid location. For this problem, goal_state=None.</p> <p>Assume N is a positive integer for row-size of the grid problem. Assume M is a positive integer for column-size of the grid problem.</p> <p>Assume wall_coords is a list of (x, y) tuple grid locations indicating the position of walls. Assume food_coords is a list of (x, y) tuple grid locations indicating the position of food.</p> <p>Create a tuple of booleans food_eaten of the same size as food_coords where food_eaten[i] indicates if food_coords[i] was eaten (true) or not (false). Initially, this tuple should be set to false (all food not eaten at start).</p> <p>Set self.initial state = (initial state, food eaten)</p>	

GridProblem Functions to Override			
Name	Arguments	Returns	implementation
actions	state	<p>Returns a list of actions that can be reached by <code>state</code>. If no action is legal, then return an empty list [].</p> <p>Note: you can use whatever format (strings, integers, etc) for the actions, but make sure to return the list of legal actions in up, down, right, left order.</p>	<p>Hint: based on walls and grid limits, you should return only legal actions. It is up to you decide on the format of the list of actions. I used strings 'up', 'down', 'right', 'left'.</p> <p>Note: you can use whatever format (strings, integers, etc) for the actions, but make sure to return the list of legal actions in up, down, right, left order.</p>
result	state, action	Returns the transition-state that results from executing action. Recall that a state is a 2-element list.	Hint: if the action is illegal, then agent should remain in <code>state</code> . If the agent moves to a grid containing food, then the transition state should reflect this through the tuple of Booleans.
action_cost	state1, action, state2	Returns the cost of transition from state1 to state2 using action.	Any action costs 1.
is_goal	state	Returns true if all food items in state consumed.	Hint: recall that a state contains a tuple of Booleans.
h	node	Returns the heuristic distance to the goal.	<p>Implement a heuristic that computes the Manhattan distance to the nearest uneaten food item.</p> <p>Note: make sure to return 0 if node is a goal-node.</p>

To clarify the usage of GridProblem, consider the following 5-by-7 (N-by-M) grid instance:



The agent is in (7, 4). The walls are at [(4,3), (5,1), (5,2)]. The food are at [(3,1), (2,3), (4,5)].

We can create this GridProblem as follows:

```
example_walls = [(4,3), (5,1), (5,2)]

example_food = [(3,1), (2,3), (4,5)]

example_grid_problem = GridProblem(initial_state=(7,4),
                                   N=5, M=7,
                                   wall_coords=example_walls,
                                   food_coords=example_food)

goal_node = breadth_first_search(example_grid_problem, tree_like=False)
```

You can try testing your search algorithms by passing this grid-problem.

Test your code on runner.py

You can test your code by running test_runner.py which will run the search algorithms on toy problems. You can inspect the code and add your own problem cases.

The following are my implementation results:

BFS					
7 generated nodes		5 popped		2 solution cost	
358 generated nodes		118 popped		11 solution cost	
365 generated nodes		123 popped		13 solution cost	
				2 solution depth	
				11 solution depth	
				13 solution depth	
				<problem.RouteProblem object a	
				<problem.GridProblem object at	
				TOTAL	

DFS					
6 generated nodes		3 popped		2 solution cost	
169 generated nodes		58 popped		26 solution cost	
175 generated nodes		61 popped		28 solution cost	
				2 solution depth	
				26 solution depth	
				28 solution depth	
				<problem.RouteProblem object a	
				<problem.GridProblem object at	
				TOTAL	

UCS					
7 generated nodes		5 popped		2 solution cost	
358 generated nodes		118 popped		11 solution cost	
365 generated nodes		123 popped		13 solution cost	
				2 solution depth	
				11 solution depth	
				13 solution depth	
				<problem.RouteProblem object a	
				<problem.GridProblem object at	
				TOTAL	

astar_treelike					
6 generated nodes		3 popped		2 solution cost	
40,276 generated nodes		12,426 popped		11 solution cost	
40,282 generated nodes		12,429 popped		13 solution cost	
				2 solution depth	
				11 solution depth	
				13 solution depth	
				<problem.RouteProblem object a	
				<problem.GridProblem object at	
				TOTAL	

astar					
6 generated nodes		3 popped		2 solution cost	
210 generated nodes		68 popped		11 solution cost	
216 generated nodes		71 popped		13 solution cost	
				2 solution depth	
				11 solution depth	
				13 solution depth	
				<problem.RouteProblem object a	
				<problem.GridProblem object at	
				TOTAL	

greedy					
6 generated nodes		3 popped		2 solution cost	
266 generated nodes		88 popped		14 solution cost	
272 generated nodes		91 popped		16 solution cost	
				2 solution depth	
				14 solution depth	
				16 solution depth	
				<problem.RouteProblem object a	
				<problem.GridProblem object at	
				TOTAL	

Notice how the treelike version of astar generated a large number of nodes, whereas graph-like astar produced much fewer and finds the same solution.

[Optional Bonus] search_algorithms.py: Visualizing solutions

Your task is to implement some helper functions in `search_algorithms.py` to help us visualize the solution paths of `RouteProblem` and `GridProblem`. We will use the [matplotlib](#) packages to ease this implementation, so add these imports to `search_algorithms.py` as follows:

```
import matplotlib.pyplot as plt
```

First let us define a function that will visualize the solution of a `RouteProblem` instance.

Name: `visualize_route_problem_solution`

Arguments: `problem`, `goal_node`, `file_name`

Returns: does not return anything.

Implementation: You may want to make use of `get_path_states` to get the states-path to `goal_node`. Using `problem` and the path of states, you can plot each state using its (x, y) coordinate which will depict the graph vertices. When plotting the states, use [scatter](#) and set `marker` to make the plot points have a square shape (see [documentation](#)). Also, be sure to set the `color` of initial state to red, the goal state to green, and the transition states to blue (see color [documentation](#)). Next, plot the possible actions (directed edges between states) as black arrows using [arrow](#). Next, plot the solution path as magenta directed arrows using [arrow](#). Finally, save the resulting figure as a png using [savefig](#). Do **not** define a figure size at the beginning of plotting, just start using `plt` to plot right away. Also, make sure at the end to call `plt.close()`.

[Optional Bonus] search_algorithms.py: Visualizing solutions

Next let us define a function that will visualize the solution of a GridProblem instance.

Name: `visualize_grid_problem_solution`

Arguments: `problem`, `goal_node`, `file_name`

Returns: does not return anything.

Implementation: You may want to make use of `get_path_states` to get the states-path to `goal_node`. Using `problem`, you can plot the location of walls, food, and agent.

- Use [scatter](#) to plot the walls as black squares. Set the size of marker to `s=2500`.
- Use [scatter](#) to plot the food as green hexagons. Set the size of marker to `s=1000`.
- Use [scatter](#) to plot the initial location of the agent as red triangle. Set the size of marker to `s=1000`.
- Plot the agent movement as magenta directed arrows using [arrow](#).
- Now, we want to adjust the limits of the plot to be more visually pleasing. So call these two commands which will set the x-axis and y-axis limits of the plot:
 - `plt.ylim([0.5, problem.N + 0.5])`
 - `plt.xlim([0.5, problem.M + 0.5])`

Finally, save the resulting figure as a png using [savefig](#). Do **not** define a figure size at the beginning of plotting, just start using `plt` to plot right away. Also, make sure at the end to call `plt.close()`.

You can test your function implementations on example `RouteProblem` and `GridProblem` to see the resulting figures.

If you run `visualize_runner.py` on your code, you can check your resulting images against the reference images: [`ref_grid.png`, `ref_route.png`]. You should reproduce the same images.